

Modern Database Systems

Lecture 2

Aristides Gionis
Michael Mathioudakis

T.A.: Orestis Kostakis



Spring 2016

logistics

- **assignment 1** is up
- **cowbook** available at learning center beta, otakaari 1 x
- if you do not have **access to the lab**, provide Aalto username or email **today!!**
- if you do not have **access at mycourses**, i will post material (slides and assignments) also at <http://michalis.co/moderndb/>

in this lecture...


b+ trees and hash-based indexing
external sorting
join algorithms
query optimization

b+ trees

b+ trees

non-leaf nodes
index entries
used to direct search

leaf nodes
contain **data-entries**
sequentially linked



each **node** stored in **one page**
data entries can be *any* one of the three *alternative* types
type 1: data records; **type 2**: (k, rid); **type 3**: (k, rids)
at least 50% capacity - **except for root!**

in the examples that follow...

alternative 2 is used


all nodes have between **d** and **2d** key entries

d is the **order** of the tree


b+ trees

non-leaf nodes
index entries
used to direct search

leaf nodes
contain **data-entries**
sequentially linked




closer look at non-leaf nodes



b+ trees

non-leaf nodes
index entries
used to direct search

leaf nodes
contain **data-entries**
sequentially linked



most widely used index


search and updates at $\log_F N$ cost (cost = pages I/O)

F = fanout (num of pointers per index node); N = num of leaf pages

efficient **equality** and **range** queries

example b+ tree - search

search begins at root, and key comparisons direct it to a leaf
search for 5*; search for all data entries $\geq 24^*$



inserting a data entry


1. find correct leaf **L**
2. place data entry onto **L**
 - a. if **L** has enough space, done!
 - b. else must split **L** into **L** and **L₂**
 - redistribute entries evenly
 - copy up the middle key to parent of **L**
 - insert entry pointing to **L₂** to parent of **L**

the above happens recursively
when index nodes are split, push up middle key

splits grow the tree
root split increases height

example b+ tree

L




insert 8*

split!


middle key is **copied up** (and continues to appear in the leaf)

L




L₂


split parent node!




middle key is **pushed up**



example b+ tree



insert 8*



deleting a data entry

inverse of insertion

deletion

remove data entry

vs

insertion

add data entry

not always used!
data tend to grow

when nodes
are less than half-full

re-distribute entries &
(maybe) merge nodes

vs

when nodes
overflow

split nodes &
re-distribute entries

b+ trees in practice

typical order $d = 100$, fill-factor = 67%
average fan-out 133

typical capacities:
for height 4: $133^4 = 312,900,700$ records
for height 3: $133^3 = 2,352,637$ records

can often hold top levels in main memory

level 1: 1 page = 8KBytes
level 2: 133 pages = 1MByte
level 3: 17,689 pages = 133 MBytes

hash-based indexes

hash-based index

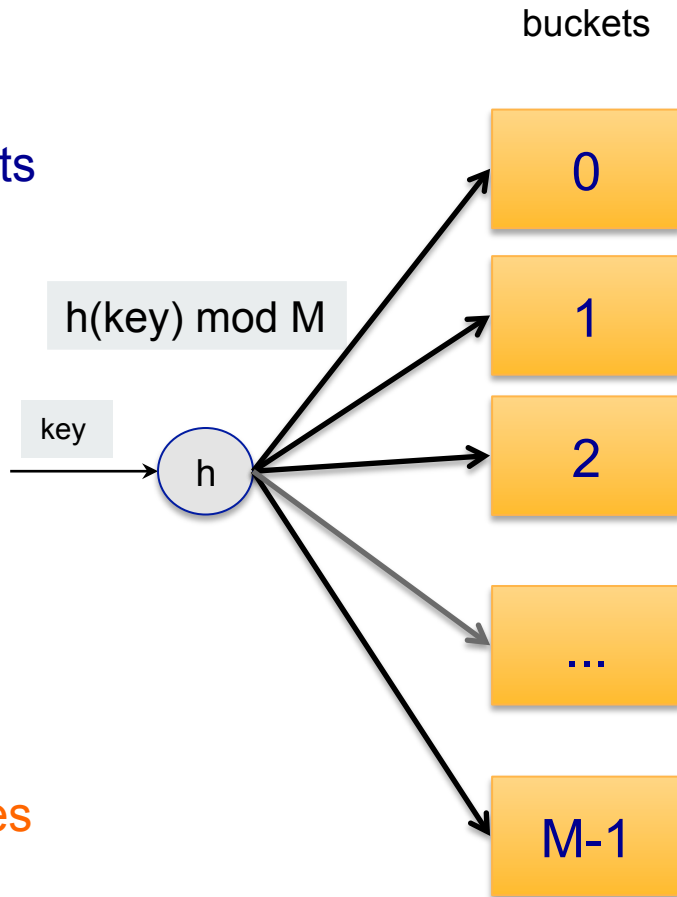
data entries organized in M buckets

bucket = a collection of pages

the data entry for record
with search key value key
is assigned to bucket
 $h(key) \bmod M$

hash function $h(key)$
e.g., $h(key) = \alpha \cdot key + \beta$

the index supports equality queries
does **not** support range queries
static and dynamic variants exist




static hashing

number of buckets is **fixed**

start with **one page** per bucket

allocated sequentially, **never de-allocated**
can use **overflow** pages



static hashing

drawback

long **overflow** chains can **degrade** performance

dynamic hashing techniques

adapt index to **data size**

extendible and *linear hashing*

extendible hashing

problem: bucket becomes full
one solution

double the number of buckets...

...and redistribute data entries

however

reading and re-writing all buckets is expensive

better idea:

use **directory of pointers** to buckets


double number of 'logical' buckets...

but split ***physically*** only the **overflown** bucket

directory much smaller than data entry pages - good!

no overflow pages - good!

example




directory is array of size $M = 4 = 2^2$
to find bucket for r , take last 2 # bits of $h(r)$
 $h(r) = \text{key}$

e.g., if $h(r) = 5 = \text{binary } 101$
it is in bucket pointed to by 01

$\text{global depth} = 2$
= min. bits enough to enumerate buckets

local depth
= min bits to identify individual bucket
= 2

insertion



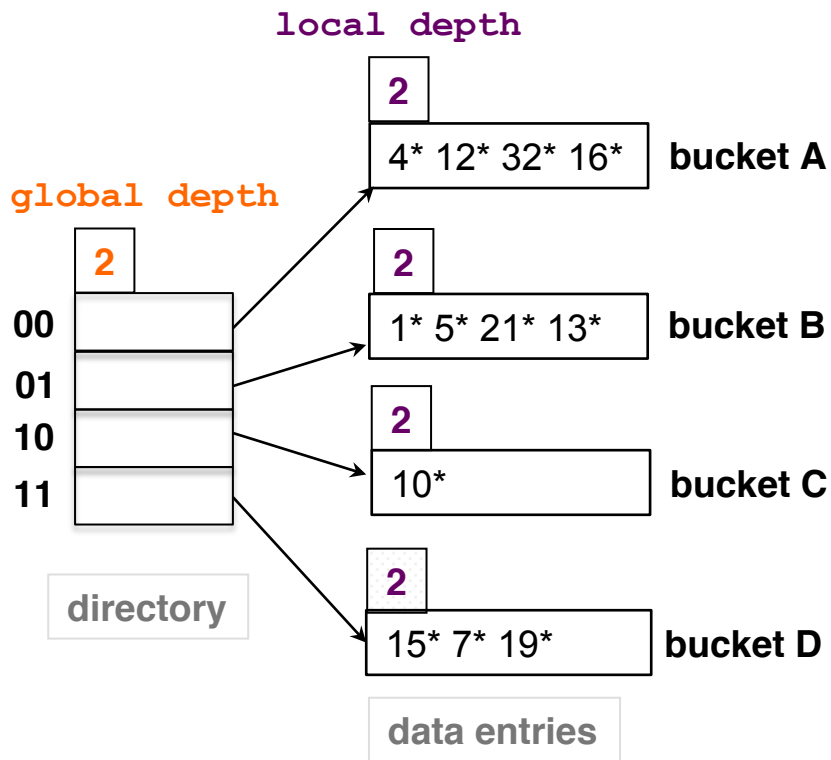
try to insert entry to
corresponding bucket

if bucket is **full**,
increase +1 **local depth**
and **split** bucket
(allocate new bucket, **re-distribute**)

if necessary, double the **directory**
i.e., when for split bucket
local depth > **global depth**

when directory doubles,
increase **global depth** +1

example



insert record with $h(r) = 20$
= binary 10100 \rightarrow bucket A

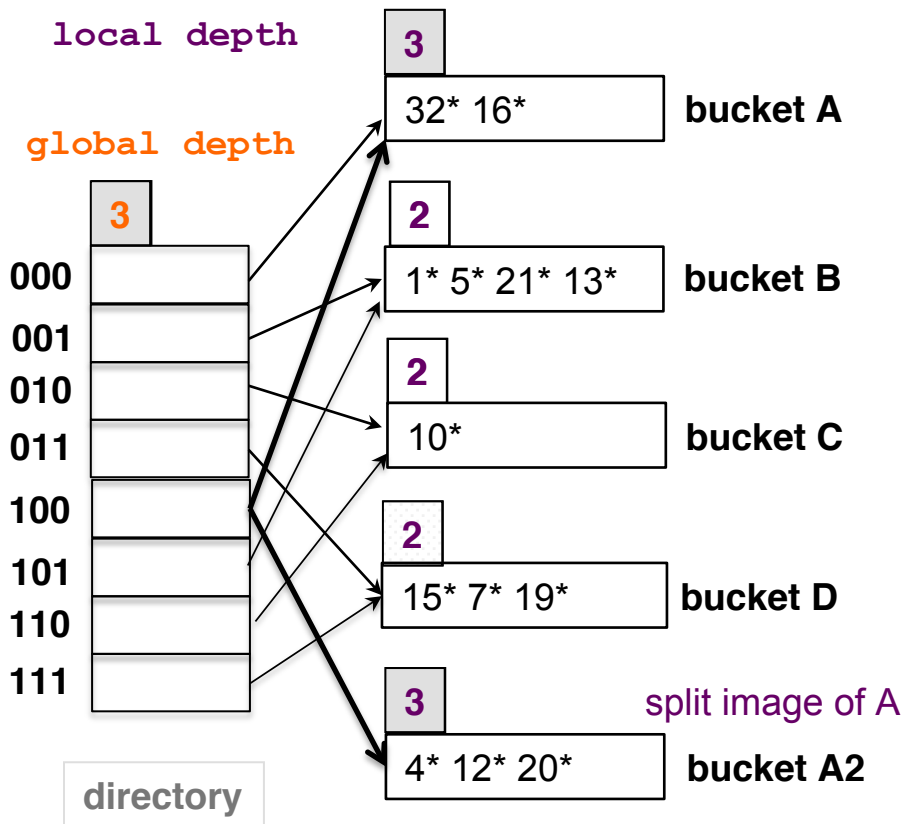
split bucket A !
allocate new page,
redistribute according to
modulo $2M = 8$
3 least significant bits

we'll have more than 4
buckets now,
so double the directory!

example

split bucket A and
redistribute entries

update local depth
double the directory
update global depth
update pointers



notes

20 = binary 10100

last 2 bits (00) tell us r belongs in A or A2

last 3 bits needed to tell which

global depth of directory

number of bits enough to determine which bucket any entry belongs to


local depth of a bucket

number of bits enough to determine if an entry belongs to this bucket

when does bucket split cause directory doubling?

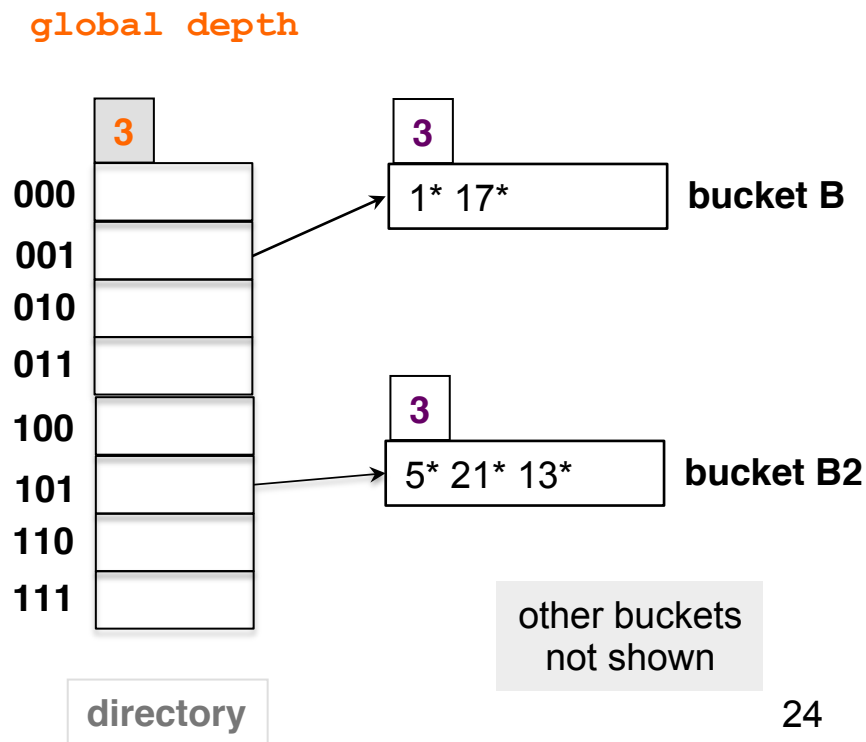
before insert, *local depth* of bucket = *global depth*

example



insert $h(r) = 17$

split bucket B



comments on extendible hashing

if directory fits in memory,
equality query answered in one disk access
answered = retrieve rid

directory grows in spurts
if hash values are skewed, it might grow large

delete: reverse algorithm
empty bucket can be merged with its 'split image'
when can the directory be halved?
when all directory elements point to same bucket as their 'split image'

indexes in SQL

create index

```
CREATE INDEX indexb  
ON students (age, grade)  
USING BTREE;
```

```
CREATE INDEX indexh  
ON students (age, grade)  
USING HASH;
```

```
DROP INDEX indexh  
ON student;
```

external sorting

the sorting problem

setting

a relation R , stored over N disk pages
 $3 \leq B < N$ pages available in memory (buffer pages)

task

sort records of R and store result on disk
sort by a function of record field values $f(r)$

why

application need records ordered
part of join implementation (soon...)


sorting with 3 buffer pages

2 phases

input
relation R


1
2
N

N pages stored on disk



buffer (memory used by dbms)

output
sorted R



N pages stored on disk

sorting with 3 buffer pages - first phase

pass 0: output N runs

run: sorted sub-file


after first phase: one run is one page
how: load one page at a time,
sort it in-memory, output to disk

output
N runs

input
relation R

N pages stored on disk

1
2
N



↓ run #1
↓ run #2
↓ run #N

only 1 buffer page needed for first phase

sorting with 3 buffer pages - second phase

pass 1,2,...: halve the runs

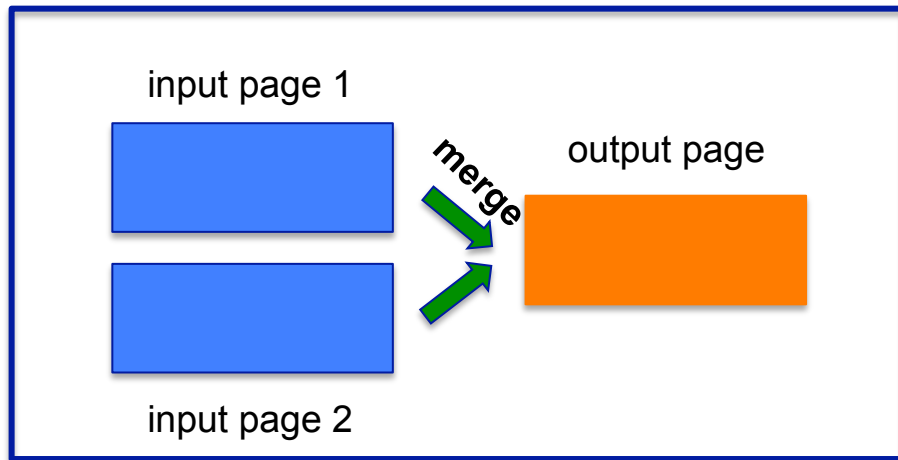
how: scan pairs of runs, each in own page,
merge in-memory into a new run,
output to disk

input
relation R

output
half runs

N pages stored on disk

run #1
run #2
run #(N-1)
run #N




run #1
...
run #N/2

N pages stored on disk

merge?

merge the two sorted input pages
into the output page
maintaining sorted order


values
 $f(r)$



compare the next smallest value from each page
move smallest to output page

merge?


merge the two input pages into the output page
maintaining sorted order



compare the next smallest value from each page
move smallest to output page

merge?


merge the two input pages into the output page
maintaining sorted order



compare the next smallest value from each page
move smallest to output page

merge?


merge the two input pages into the output page
maintaining sorted order




compare the next smallest value from each page
move smallest to output page

merge?

merge the two input pages into the output page
maintaining sorted order



merge




output page is full,
what do we do?

write it to disk!

compare the next smallest value from each page
move smallest to output page

merge?


merge the two input pages into the output page
maintaining sorted order



compare the next smallest value from each page
move smallest to output page

merge?


merge the two input pages into the output page
maintaining sorted order



compare the next smallest value from each page
move smallest to output page

merge?

merge the two input pages into the output page
maintaining sorted order




compare the next smallest value from each page
move smallest to output page

merge?

merge the two input pages into the output page
maintaining sorted order

if the input run has more pages, load next one




input page is empty!
what do we do?

next smallest value from each page
move smallest to output page

merge?


merge the two input pages into the output page
maintaining sorted order



compare the next smallest value from each page
move smallest to output page

merge?

merge the two input pages into the output page
maintaining sorted order



compare the next smallest value from each page
move smallest to output page

sorting with 3 buffer pages - second phase


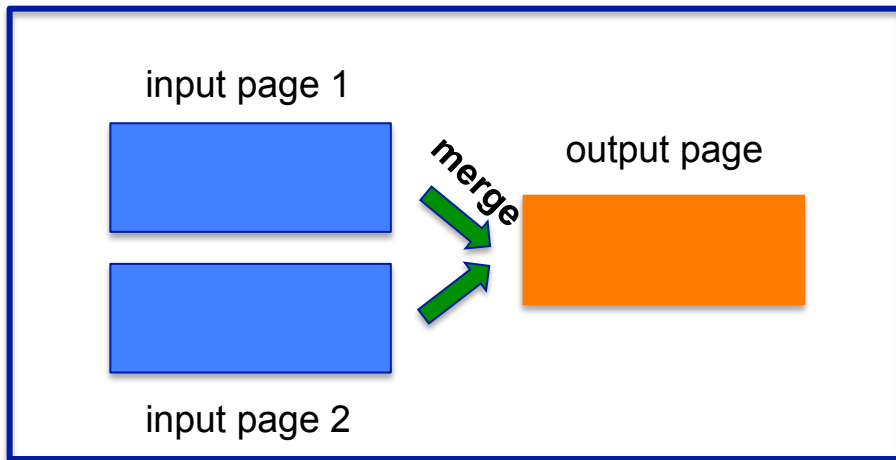
pass 1,2,...: halve the runs
after $\log_2 N$ passes...
we are done!

input
relation R

output
sorted R

N pages stored on disk

run #1
run #2
run #(N-1)
run #N



N pages stored on disk

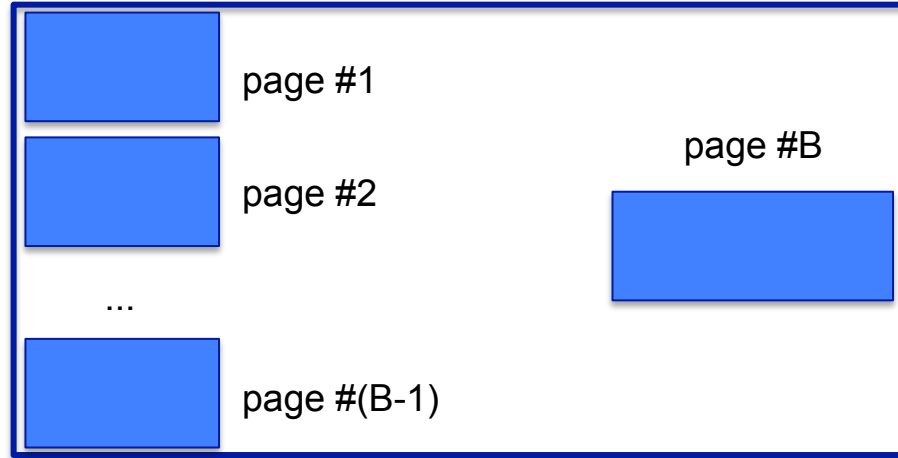
sorting with B buffer pages

same approach


input
relation R

1
2
N

N pages stored on disk



output
sorted R



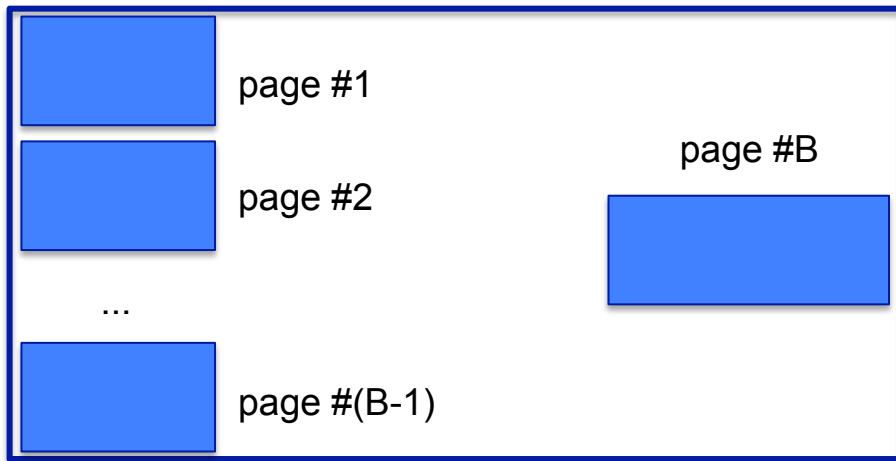
N pages stored on disk

sorting with B buffer pages - first phase

pass 0: output $\lceil N/B \rceil$ runs
how: load R to memory in chunks
of B pages, sort in-memory,
output to disk

input
relation R

1
2
N



output
sorted R

run #1
run #2
run $\lceil N/B \rceil$

N pages stored on disk

N pages stored on disk

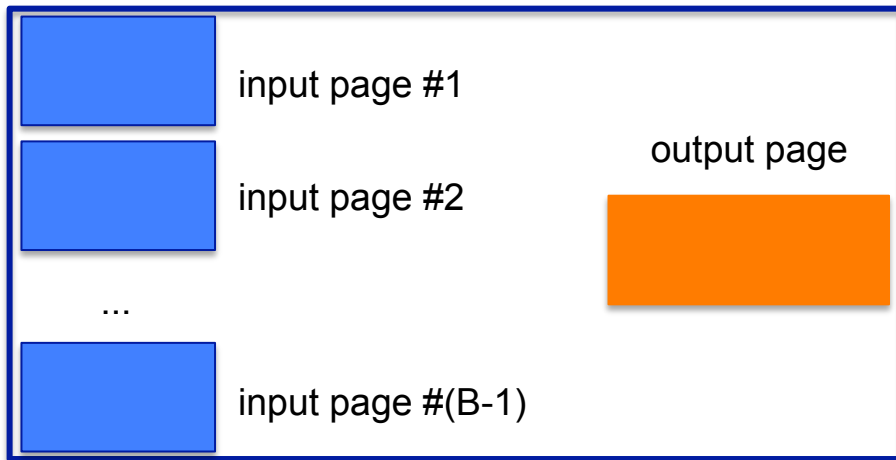
sorting with B buffer pages - second phase

pass 1,2,...:
merge runs in groups of B-1


input
relation R

run #1
run #2
run #[N/B]

N pages stored on disk



output
sorted R



N pages stored on disk

sorting with B buffer pages

how many passes in total?


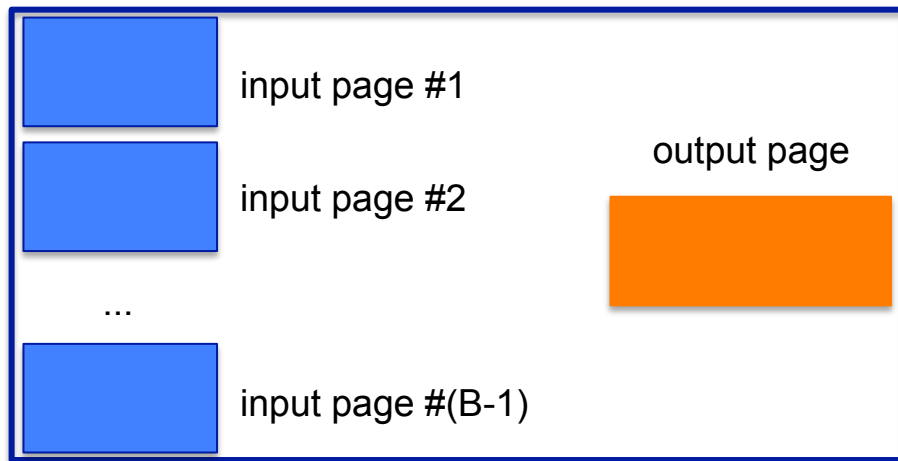
$$\text{let } N_1 = \lfloor N/B \rfloor$$

total number of passes =

input relation R phase 1 \longrightarrow $1 + \lceil \log_{B-1}(N_1) \rceil$ \longleftarrow phase 2 output sorted R

N pages stored on disk

1
2
N



N pages stored on disk

sorting with B buffer

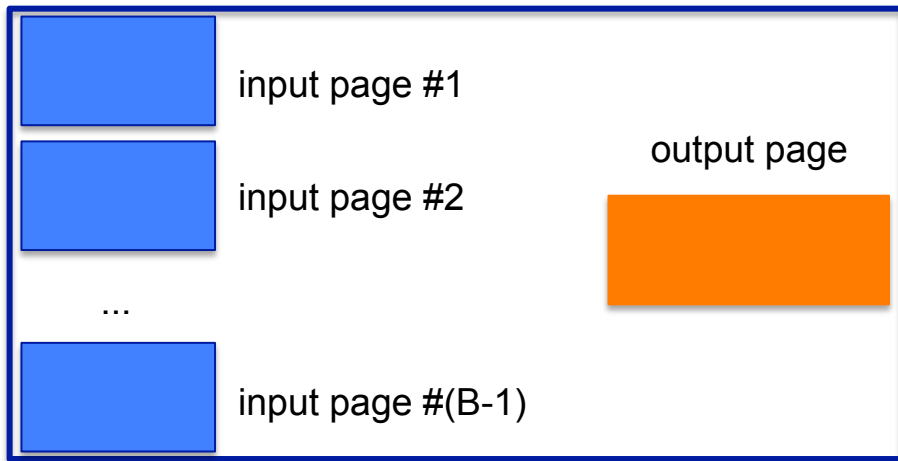
how many pages I/O per pass?

$2N$: N input, N output


input
relation R

1
2
N

N pages stored on disk



output
sorted R



N pages stored on disk

sql joins

joins

so far, we have seen **queries** that
operate on a **single relation**

but we can also have queries that
combine information
from **two or more relations**

joins

students

sid	name	username	age
53666	Sam Jones	jones	22
53688	Alice Smith	smith	22
53650	Jon Edwards	jon	23

dbcourse

sid	points	grade
53666	92	A
53650	65	C

what does this compute?

```
SELECT *  
FROM students S, dbcourse C  
WHERE S.sid = C.sid
```

S.sid	S.name	S.username	C.age	C.sid	C.points	C.grade
53666	Sam Jones	jones	22	53666	92	A
53650	Jon Edwards	jon	23	53650	65	C

joins

```
SELECT *  
FROM students S, dbcourse C  
WHERE S.sid = C.sid
```

intuitively...
take all pairs of records from S and C
(the “cross product” $S \times C$)

keep only records that
satisfy WHERE condition

S record #1	C record #1
S record #1	C record #2
S record #1	C record #3
...	...
S record #2	C record #1
S record #2	C record #2
S record #2	C record #3
...	...

joins

```
SELECT *  
FROM students S, dbcourse C  
WHERE S.sid = C.sid
```

intuitively...

take all pairs of records from S and C
(the “cross product” S x C)

keep only records that
satisfy WHERE condition

output join result

$S \bowtie_{S.sid=C.sid} C$

S record #1	C record #1
S record #1	C record #2
S record #1	C record #3
...	...
S record #2	C record #1
S record #2	C record #2
S record #2	C record #3
...	...

joins

expensive to
materialize!

```
SELECT *  
FROM students S, dbcourse C  
WHERE S.sid = C.sid
```

intuitively...

take all pairs of records from S and C
(the “cross product” S x C)

keep only records that
satisfy WHERE condition

output join result

$S \bowtie_{S.sid=C.sid} C$

S record #1	C record #2
S record #2	C record #1

in what follows...

```
SELECT *  
FROM students S, dbcourse C  
WHERE S.sid = C.sid
```

algorithms to compute joins
without materializing cross product

assuming WHERE condition is equality condition
as in the example
assumption is not essential, though

join algorithms

the join problem

input

relation **R**: M pages on disk, p_R records per page

relation **S**: N pages on disk, p_S records per page

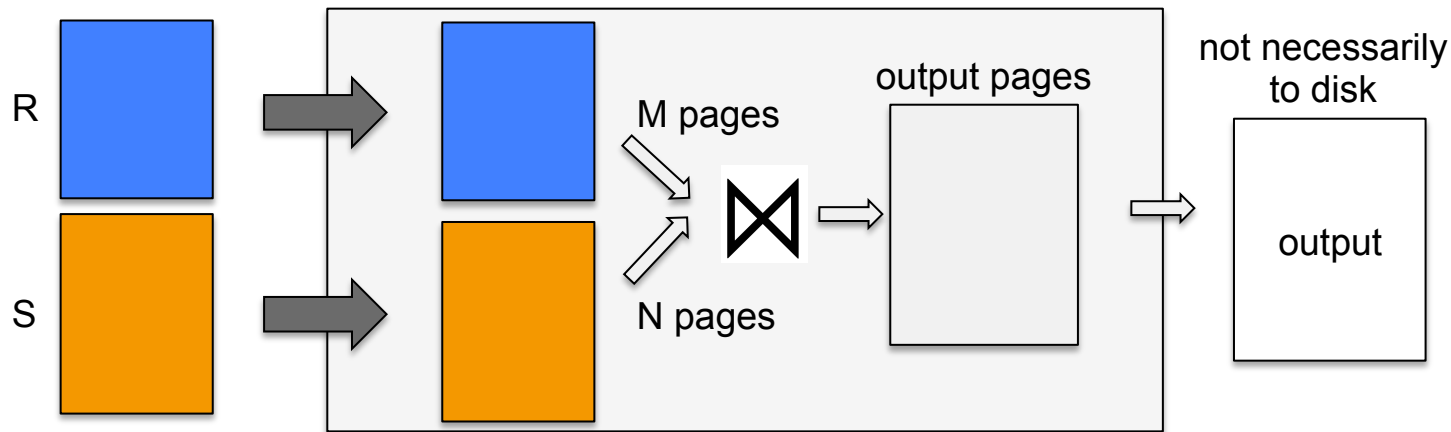
$$M \leq N$$

output

$$R \bowtie_{R.a=S.b} S$$

if there is enough memory...

load both relations in memory



we only have to scan
each relation once


in-memory

for each record \mathbf{r} in R
for each record \mathbf{s} in S
if $\mathbf{r.a} = \mathbf{s.b}$:
store (\mathbf{r}, \mathbf{s}) in output pages
output

I/O cost
(ignoring final output cost)
M + N pages

page-oriented simple nested loops join

join using **3** memory (buffer) pages




R is scanned once
S is scanned M times

for each page P of R
for each page Q of S
compute P join Q; store in output page

I/O cost (pages)
 $M + M * N$

block nested loops join

join using **B** memory (buffer) pages




R is scanned once
S is scanned $M/(B-2)$ times

for each **block** P of **(B-2)** pages of R
for each page Q of S
compute P join Q; store in output page

I/O cost (pages)
 $M + [M/(B-2)] * N$ ₆₁

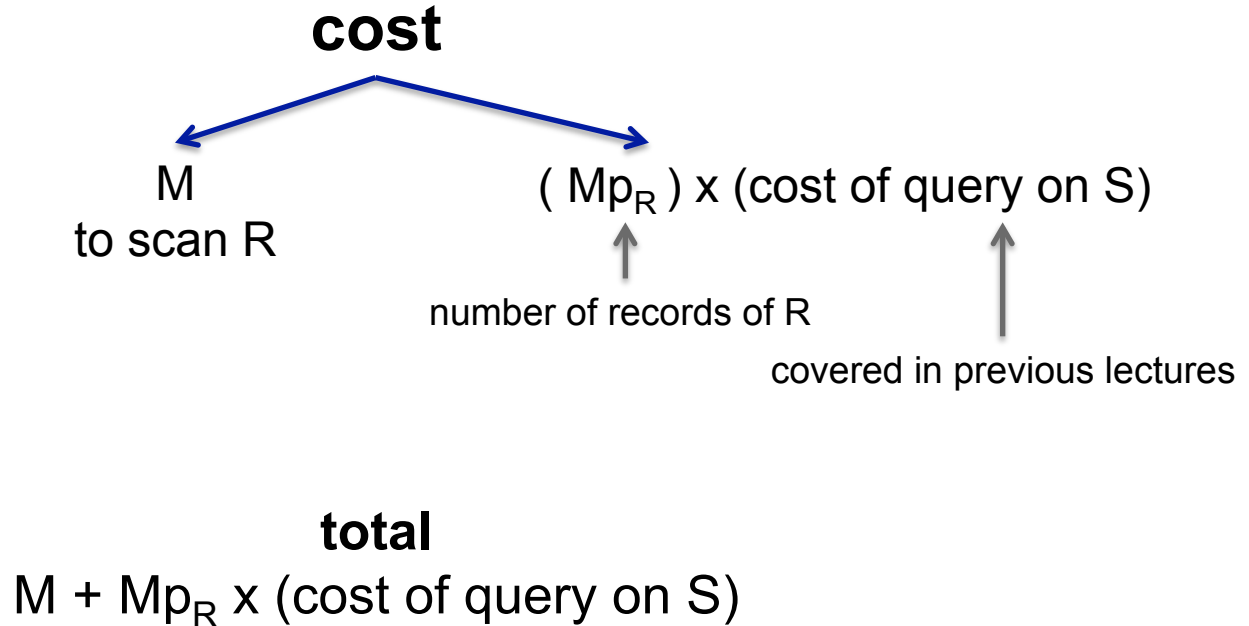
index nested loop join

relation S has an index on the join attribute
use one page to make a pass over R
use index to retrieve only matching records of S



```
for each record r of R
  for each record s of S with s.b = r.a // query index
    add (r,s) to output
```

index nested loop join



sort-merge join

two phases

sort and merge

sort

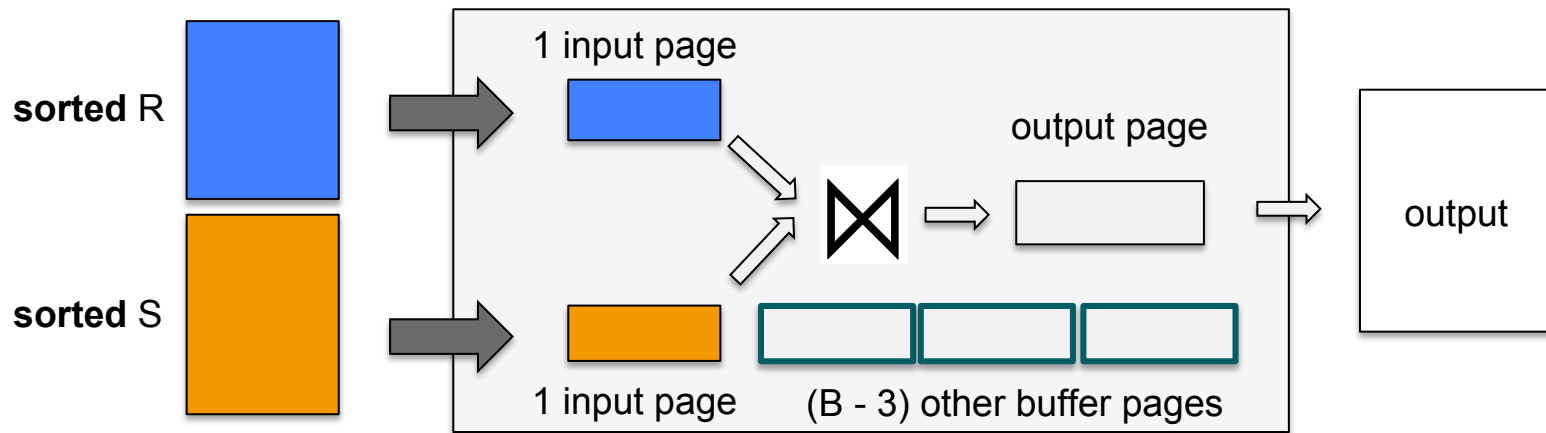
R and S on the join attribute

using **external sort algorithm**

cost $O(n \log n)$, n : number of relation pages

merge sorted R and S

sort-merge join: merge



sort-merge join: merge

assumption
b is a key for S

	R.a	S.b	
r →	1	3	← S
	2	4	
	6	5	
	7	6	
	8	9	
	9	10	

current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

r →	1	3 ← S
	2	4
	6	5
	7	6
	8	9
	9	10

current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
    advance r until r >= s  
    advance s until s >= r  
    output (r, s)  
    advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

	1	3	← S
	2	4	
r →	6	5	
	7	6	
	8	9	
	9	10	

current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
    advance r until r >= s  
    advance s until s >= r  
    output (r, s)  
    advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10



current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10



current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10



current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10



main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

current pages
in memory
(only join attributes
are shown)

sort-merge join: merge

assumption
b is a key for S

current pages
in memory
(only join attributes
are shown)

R.a	S.b
1	3
2	4
6	5
7	6
8	9
9	10

r →

← s

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10

r →

← s

current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
    advance r until r >= s  
    advance s until s >= r  
    output (r, s)  
    advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10

r →

← s

current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

R.a **S.b**

1	3
2	4
6	5
7	6
8	9
9	10



current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption

b is a key for S

new page
for R!

r →

R.a

S.b

13

3

14

4

15

5

17

6

19

9

23

10

← s

current pages
in memory
(only join attributes
are shown)

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

sort-merge join: merge

assumption
b is a key for S

current pages
in memory
(only join attributes
are shown)

	R.a	S.b
r →	13	3
	14	4
	15	5
	17	6
	19	9
	23	10 ← s

main loop

```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

cost for merge
 $M + N$

sort-merge join: merge

what if this assumption
does not hold?

assumption
b is a key for S

current pages
in memory
(only join attributes
are shown)

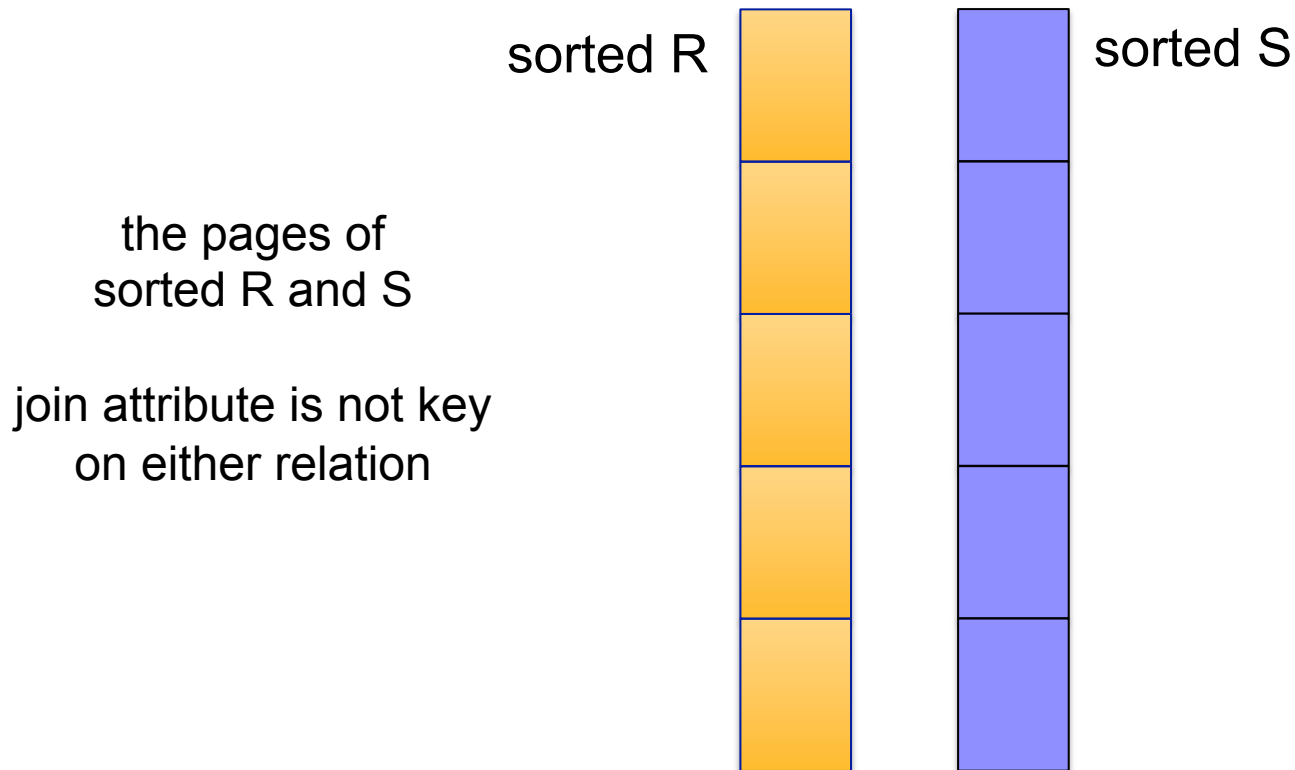
	R.a	S.b
r →	13	3
	14	4
	15	5
	17	6
	19	9
	23	10 ← s

main loop

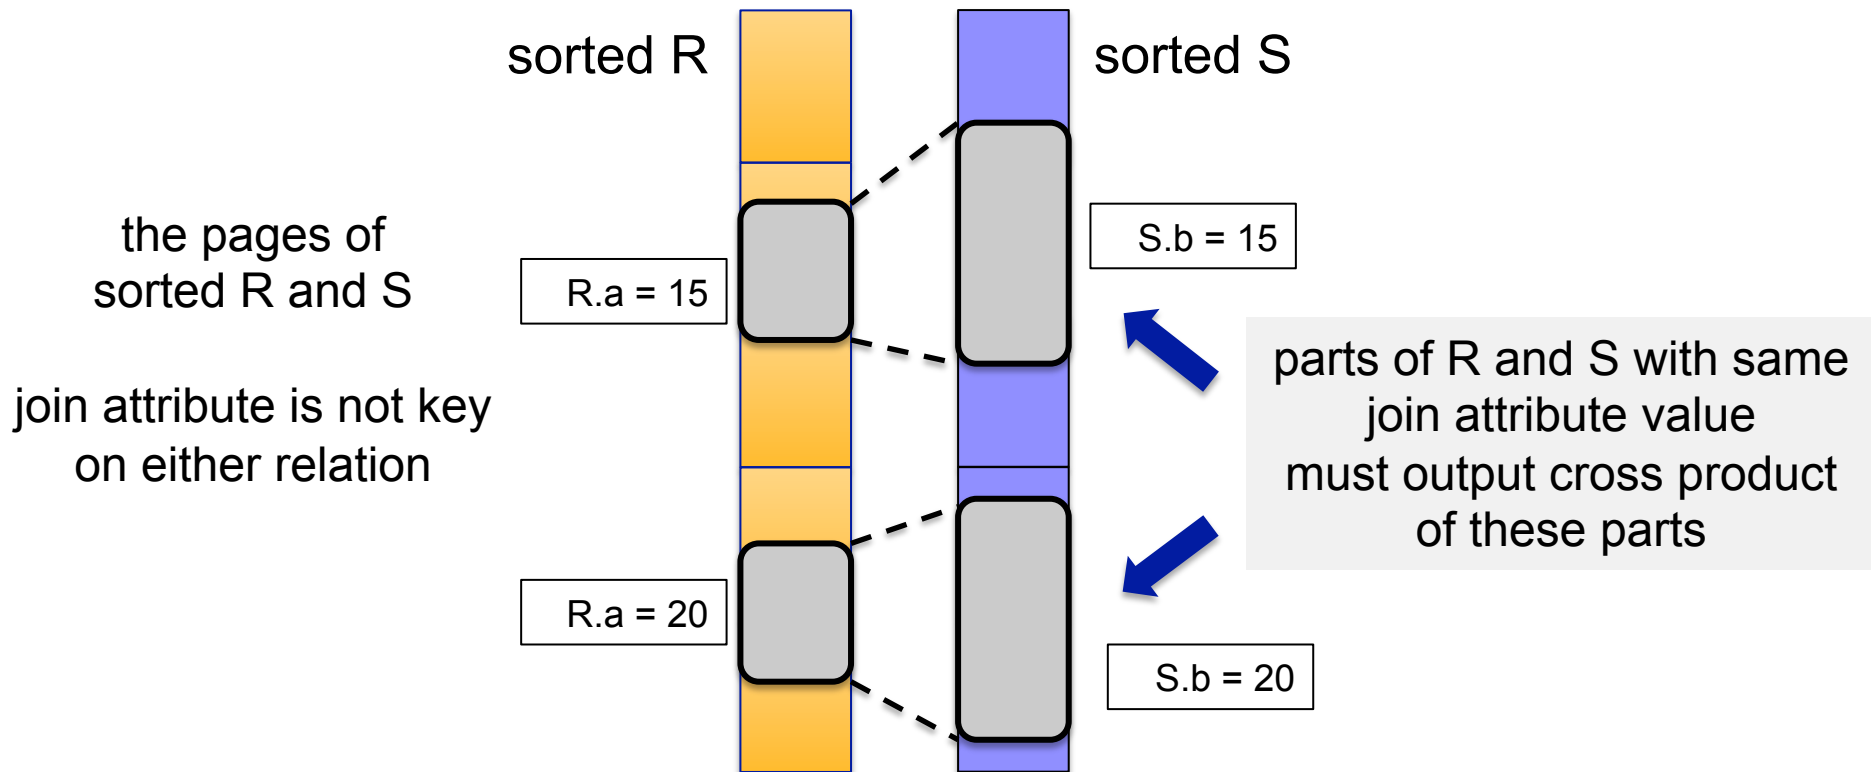
```
repeat while r != s:  
  advance r until r >= s  
  advance s until s >= r  
  output (r, s)  
  advance r and s
```

cost for merge
 $M + N$

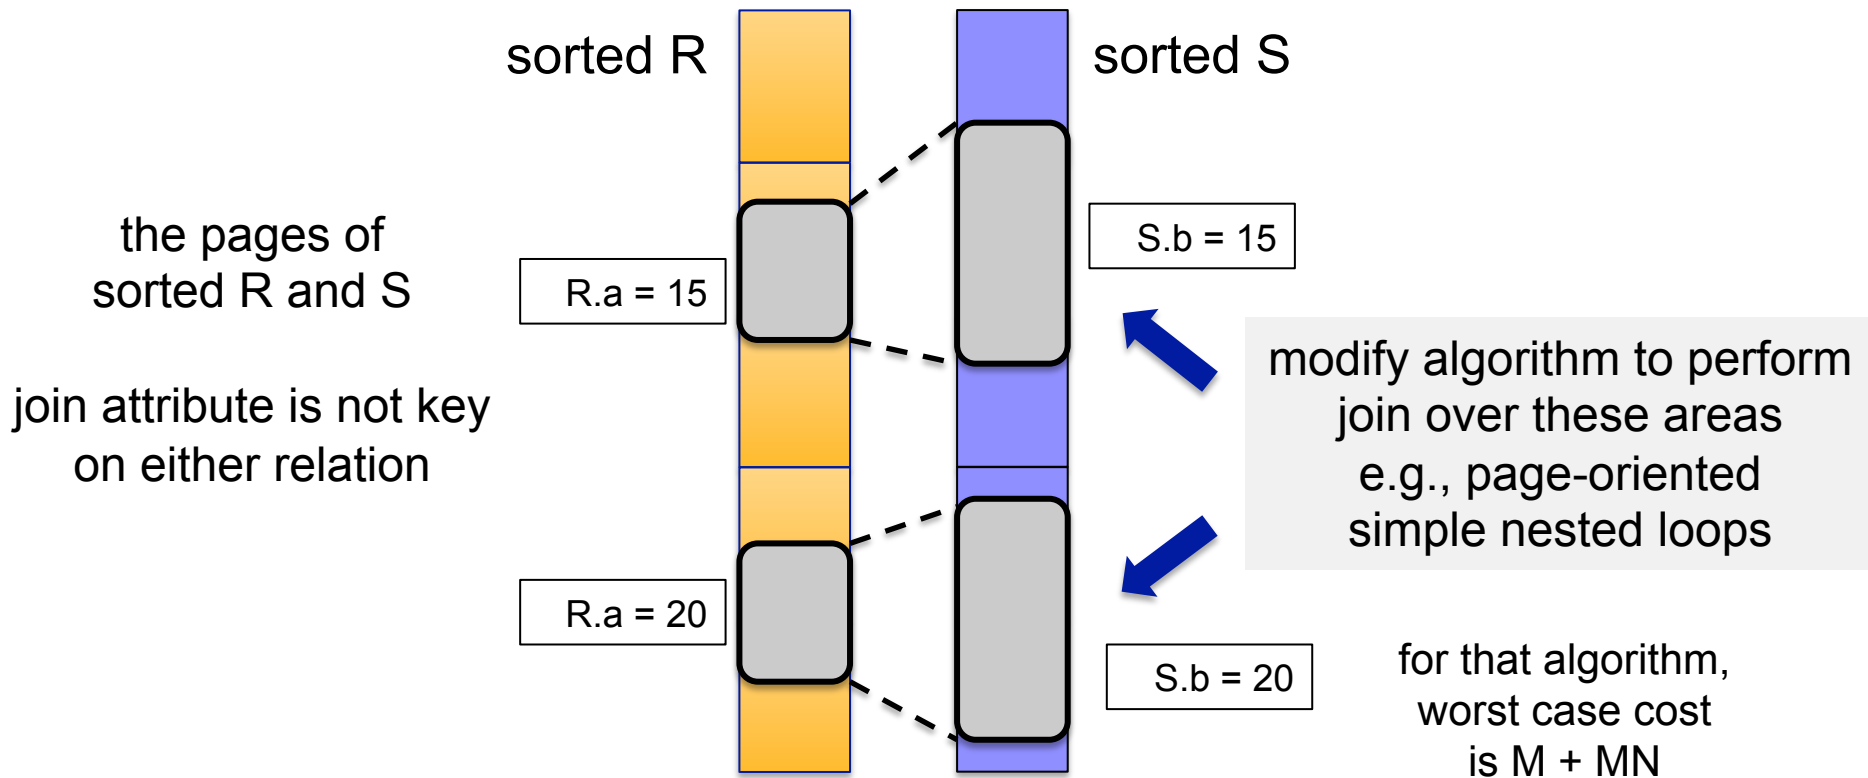
sort-merge join: merge



sort-merge join: merge



sort-merge join: merge



hash join

two phases

partition and probe

partition


each relation into partitions
using the **same hash function**
on the **join attribute**

probe

join the **corresponding** partitions

hash join - partition


B buffer pages available
scan R with 1 buffer page
hash into B-1 partitions
on **join attribute**



use B-1 pages to hold the partitions,
flush when full or scan of R ends

hash join - partition

B buffer pages available
scan S with 1 buffer page
hash into B-1 partitions
on **join attribute**



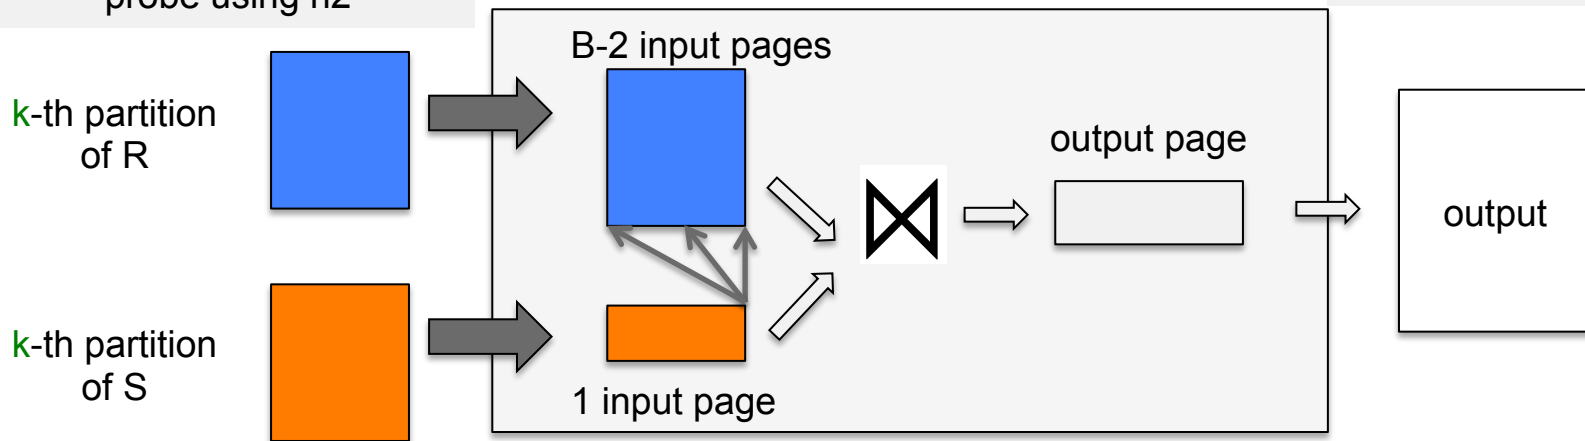
use B-1 pages to hold the partitions,
flush when full or scan of S ends

hash join - probe

variant
re-partition the partition of
R in-memory with hash
function h_2 ,
probe using h_2

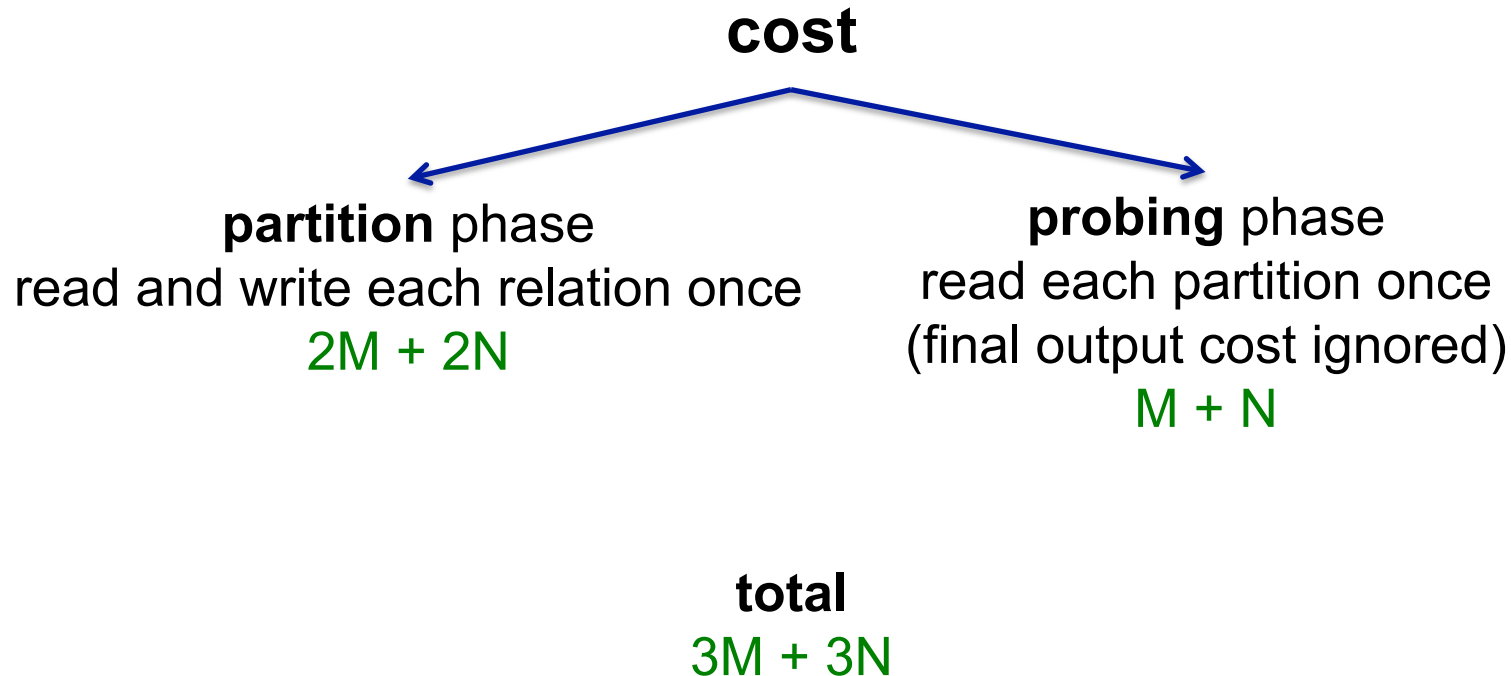
B buffer pages available
load k -th partition of R into memory
(assuming it fits in $B-2$ pages)

holds when size of each
partition fits in $B-2$ pages
approximately
 $B-2 > M / (B - 1)$
 $B > \sqrt{M}$



scan k -th partition of S one page at a time;
for each record of S, probe the partition
of R for matching records;
store matches in output page;
flush when full or done

hash join



a few words on
query optimization

query optimization

once we submit a query
the dbms is responsible for
efficient computation

the same query can be
executed in many ways
each is an '**execution plan**'
or 'query evaluation plan'

example

```
select *  
from students  
where sid = 100
```


execution plan

annotated relational
algebra tree

$\sigma_{\text{sid} = 100}$ (on-the-fly)  algorithm used by operator



students (scan)

 **access path**
how we retrieve data from
the relation: scan or index?

another execution plan

$\sigma_{\text{sid} = 100}$ (query index stud_btree)




students (b+ tree index stud_btree on sid)

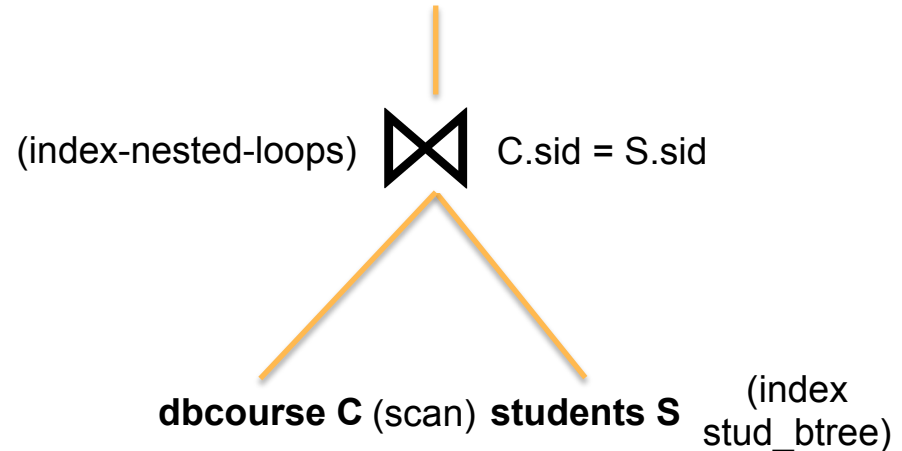
example

```
select *  
from students S, dbcourse C  
where S.sid = C.sid
```

execution plan



another execution plan



which plan to choose?

dbms **estimates cost** for a number of **execution plans**
(not all possible plans, necessarily!)

the estimates follow the cost analysis
we presented earlier

dbms picks the execution plan with
minimum estimated cost

summary

summary

- commonly used indexes
- B+ tree
 - most commonly used
 - supports efficient equation and range queries
- hash-based indexes
 - extendible hashing uses directory, not overflow pages
- external sorting
- joins
- query optimization

tutorial

next week

references

- “cowbook”, database management systems, by ramakrishnan and gehrke
- “elmasri”, fundamentals of database systems, elmasri and navathe
- other database textbooks

credits

some slides based on material from
database management systems, by ramakrishnan and gehrke

backup slides


b+ tree - deletion

deleting a data entry


1. start at root, find leaf L of entry
2. remove the entry, if it exists
 - if L is at least half-full, done!
 - else
 - try to re-distribute, borrowing from sibling
 - adjacent node with same parent as L
 - if that fails, merge L into sibling
 - if merge occurred,
 - must delete L from parent of L

merge could propagate to root


example b+ tree




delete 19*



example b+ tree




delete 20*




occupancy below 50%, redistribute!

example b+ tree




delete 20*




occupancy below 50%, redistribute!

example b+ tree




delete 24*




occupancy below 50%, merge!

example b+ tree




delete 24*




delete from parent!
reverse of copying up

occupancy below 50%, merge!


example b+ tree




delete 24*



example b+ tree




delete 24*




delete from parent!
reverse of copying up


example b+ tree




delete 24*




example b+ tree




delete 24*



example b+ tree



delete 24*




example b+ tree

during deletion of 24* -- different example

can redistribute entries of index nodes
pushing through root splitting entry

left child of root has **many entries** (full)




example b+ tree


during deletion of 24^* -- different example

can redistribute entries of index nodes
pushing through root splitting entry

left child of root has **many entries** (full)



example b+ tree



linear hashing

linear hashing

dynamic hashing

uses **overflow** pages; **no directory**
splits a bucket in **round-robin** fashion
when an **overflow** occurs

$M = 2^{\text{level}}$: number of buckets at beginning of round
pointer **next** $\in [0, M)$ points at **next bucket to split**
already **next - 1** '**split-image**' buckets **appended** to **original M**

linear hashing

to allocate entries, use

$$H_0(\text{key}) = h(\text{key}) \bmod M, \text{ or}$$

$$H_1(\text{key}) = h(\text{key}) \bmod 2M$$

i.e., **level** or **level+1** least significant bits of $h(\text{key})$

to allocate bucket for **key**

first use $H_0(\text{key})$


if $H_0(\text{key})$ is less than **next**

then it refers to a **split bucket**

use $H_1(\text{key})$ to determine if it refers
to **original** or its **split image**

linear hashing

in the middle of a round...



linear hashing inserts

insert

find bucket by applying H_0 / H_1 and
insert **if there is space**

if bucket to insert into is **full**:
add **overflow** page, insert data entry,
split next bucket and **increment next**

since buckets are split round-robin,
long overflow chains don't develop!

example - insert $h(r) = 43$

on split, H_1 is used to **redistribute** entries

