

Practical 3

Implementation of B+ Tree

GAHAN M. SARAIYA, 18MCEC10

18mcec10@nirmauni.ac.in

I. INTRODUCTION

Aim of this practical is to implement algorithm of B+ tree.

Supported Operation

- Insert single item
- Insertion in bulk
- Deletion
- Search
- Range Search

II. IMPLEMENTATION

```
1  #!/usr/bin/python3
2  # coding=utf-8
3  """
4  By Gahan Saraiya
5  B+ Tree Implementation :
   → https://github.com/gahan9/DS\_lab/blob/master/btree\_implementation/bPlusTree.py
6
7  some terminology of python used for variables in this code:
8  _var => (convention only) underscore prefix is just a hint to programmer that a
   → variable or method starting with a single underscore is intended for internal
   → use
9  var_ => (convention only) to brake name conflict
10 __var => "dunders" (name mangling) rewrite the attribute name in order to avoid
   → naming conflicts in subclasses.
11     interpreter changes the name of the variable in a way that makes it
   → harder to create collisions when the class is extended later.
12 __var__ => magic methods
13 // => Integer division (5//2 returns 2)
14 """
15
16 __author__ = "Gahan Saraiya"
17
```

```

18 __link__ =
    ↪ "https://github.com/gahan9/DS_lab/blob/master/btree_implementation/bPlusTree.py"
19
20 import math
21 import logging
22 import os
23 import random
24 from datetime import datetime
25 from bisect import bisect_right, bisect_left
26 from collections import deque
27
28 DEBUG = False
29 LOG_DIR = "."
30 logger = logging.getLogger('bPlusTree')
31 logging.basicConfig(level=logging.DEBUG,
32                     format='%(asctime)s [%(name)-8s - %(levelname)s]:
33                     ↪ %(message)s',
34                     datefmt='[%Y-%d-%m_%H.%M.%S]',
35                     filename=os.path.join(LOG_DIR, 'b_plus_tree.log'),
36                     filemode='w')
37
38 ch = logging.StreamHandler() # create console handler with a higher log level
39 ch.setLevel(logging.DEBUG)
40 # create formatter and add it to the handlers
41 formatter = logging.Formatter('%(asctime)s [%(name)-8s - %(levelname)s]:
42 ↪ %(message)s')
43 ch.setFormatter(formatter)
44 # add the handlers to the logger
45 logger.addHandler(ch)
46
47 def log(*msg):
48     if DEBUG:
49         logger.debug(msg)
50     else:
51         pass
52
53 class InternalNode(object):
54     """
55     Class : B+ Tree Internal Node
56     represents internal (non-leaf) node in B+ tree
57     """
58     def __init__(self, degree=4):
59         """
60         initialize B tree node
61         :param degree: specify degree of btree # default degree set to 4
62         """
63         self.degree = degree

```

```

63     self.keys = [] # store keys/data values
64     self.children = [] # store child nodes (list of instances of BTreeNode);
        ↳ empty list if node is leaf node
65     self.parent = None
66
67     def __repr__(self):
68         return " | ".join(map(str, self.keys))
69
70     @property
71     def is_leaf(self):
72         return False
73
74     @property
75     def total_keys(self):
76         return len(self.keys)
77
78     @property
79     def is_balanced(self):
80         # return False if total keys exceeds max accommodated keys (degree - 1)
81         return self.total_keys <= self.degree - 1
82
83     @property
84     def is_full(self):
85         return self.total_keys >= self.degree
86
87     @property
88     def is_empty(self):
89         return self.total_keys < (self.degree + 1) // 2
90
91
92     class LeafNode(object):
93         """
94         Class : B+ Tree Leaf Node
95         represents leaf node in B+ tree
96         """
97         def __init__(self, degree=4):
98             self.degree = degree
99             self.keys = [] # data values
100             self.sibling = None # sibling node to point
101             self.parent = None # parent node - None for root node
102
103         def __repr__(self):
104             return " | ".join(map(str, self.keys))
105
106         @property
107         def is_leaf(self):
108             return True
109

```

```

110     @property
111     def total_keys(self):
112         return len(self.keys)
113
114     @property
115     def is_balanced(self):
116         # return False if total keys exceeds max accommodated data (degree - 1)
117         return self.total_keys < self.degree
118
119     @property
120     def is_full(self):
121         return not self.is_balanced
122
123     @property
124     def is_empty(self):
125         return self.total_keys < math.floor(self.degree / 2)
126
127
128     class BPlusTree(object):
129         def __init__(self, degree=4):
130             self.degree = degree
131             self.__root = LeafNode(degree=degree)
132             self.__leaf = self.__root
133
134         def search_key(self, start_node, value):
135             """
136
137             :param start_node: get root or any non leaf node
138             :param value: value to be search
139             :return: most matching leaf node
140             """
141             if start_node.is_leaf:
142                 _index = bisect_left(start_node.keys, value)
143                 return _index, start_node
144             else:
145                 _index = bisect_right(start_node.keys, value)
146                 return self.search_key(start_node.children[_index], value)
147
148         def search(self, start=None, end=None):
149             """
150
151             :param start: specify start node to search range for
152             :param end: specify end node for range search
153             :return:
154             """
155             _result = []
156             node = self.__root
157             leaf = self.__leaf

```

```
158
159     if start is None:
160         while True:
161             for value in leaf.keys:
162                 if value <= end:
163                     _result.append(value)
164                 else:
165                     return _result
166             if leaf.sibling is None:
167                 return _result
168             else:
169                 leaf = leaf.sibling
170     elif end is None:
171         _index, leaf = self.search_key(node, start)
172         _result.extend(leaf.keys[_index:]) # equivalent to _result + leaf
173         while True:
174             if leaf.sibling is None:
175                 return _result
176             else:
177                 leaf = leaf.sibling
178                 _result.extend(leaf.keys)
179     else:
180         if start == end:
181             _index, _node = self.search_key(node, start)
182             try:
183                 if _node.keys[_index] == start:
184                     _result.append(_node.keys[_index])
185                     return _result
186                 else:
187                     return _result
188             except IndexError:
189                 return _result
190         else:
191             _index1, _node1 = self.search_key(node, start)
192             _index2, _node2 = self.search_key(node, end)
193             if _node1 is _node2:
194                 if _index1 == _index2:
195                     return _result
196                 else:
197                     _result.extend(_node1.keys[_index1:_index2])
198                     return _result
199             else:
200                 _result.extend(_node1.keys[_index1:])
201                 node_ = _node1
202                 while True:
203                     if _node1.sibling == _node2:
204                         _result.extend(_node2.keys[:_index2 + 1])
205                         return _result
```

```

206         else:
207             try:
208                 _result.extend(node_.sibling.keys)
209                 node_ = node_.sibling
210             except AttributeError:
211                 return _result
212
213     def traverse(self, _node):
214         _result = []
215         _result.extend(_node.keys)
216         if getattr(_node, "sibling", None) is None:
217             return _result
218         for i in range(0, len(_node.sibling))[:-1]:
219             _result.extend(self.traverse(_node.sibling[i]))
220         while True:
221             pass
222
223     def pretty_print(self):
224         # print("B+ Tree:")
225         queue, height = deque(), 0
226         queue.append([self.__root, height])
227         while True:
228             try:
229                 node, height_ = queue.popleft()
230                 # print("adding node: {}".format(node))
231             except IndexError:
232                 return
233             else:
234                 if not node.is_leaf:
235                     print("Internal Node : {:} \theight >> {}".format(node.keys,
236                                     ↪ height_))
237                     if height_ == height:
238                         height += 1
239                     queue.extend([[i, height] for i in node.children])
240                 else:
241                     print("Leaf Node      : {} \theight >> {}".format([i for i in
242                                     ↪ node.keys], height_))
243
244     def insert(self, value):
245         # log("parent:{} leaf:{} node:{}\tkeys:{}\t
246         ↪ children:{}".format(node.parent, node.is_leaf, node, node.keys,
247         ↪ getattr(node, 'children', '0')))
248
249     def split_leaf_node(node):
250         log("splitting leaf node: {}".format(node.keys))
251         mid = self.degree // 2 # integer division in python3
252         new_leaf = LeafNode(self.degree)
253         new_leaf.keys = node.keys[mid:]

```

```

250     if node.parent is None:  # None and 0 are to be treated as different
251         ↪ value
252         parent_node = InternalNode(self.degree)  # create new parent for
253         ↪ node
254         parent_node.keys, parent_node.children = [node.keys[mid]], [node,
255         ↪ new_leaf]
256         node.parent = new_leaf.parent = parent_node
257         self.__root = parent_node
258     else:
259         _index = node.parent.children.index(node)
260         node.parent.keys.insert(_index, node.keys[mid])
261         node.parent.children.insert(_index + 1, new_leaf)
262         new_leaf.parent = node.parent
263         if not node.parent.is_balanced:
264             split_internal_node(node.parent)
265     node.keys = node.keys[:mid]
266     node.sibling = new_leaf
267     log("{} --- {} --- {}".format(node, node.sibling,
268     ↪ self.__root.children))
269
270 def split_internal_node(node_):
271     mid = self.degree // 2  # integer division in python3
272     new_node = InternalNode(self.degree)
273     new_node.keys = node_.keys[mid:]
274     new_node.children = node_.children[mid:]
275     new_node.parent = node_.parent
276     for child in new_node.children:
277         child.parent = new_node  # assign parent to every new child of
278         ↪ current node
279     if node_.parent is None:  # again Note that None and 0 are not same
280         ↪ but both treated as False in boolean
281         # need to make new root if we are to split root node
282         new_root = InternalNode(self.degree)
283         new_root.keys = [node_.keys[mid - 1]]
284         new_root.children = [node_, new_node]
285         node_.parent = new_node.parent = new_root  # set parent of newly
286         ↪ created node
287         self.__root = new_root  # set new ROOT node
288     else:
289         # if node is not root internal node
290         _index = node_.parent.children.index(node_)
291         node_.parent.keys.insert(_index, node_.keys[mid - 1])
292         node_.parent.children.insert(_index + 1, new_node)
293         if not node_.parent.is_balanced:
294             split_internal_node(node_.parent)
295     node_.keys = node_.keys[:mid - 1]
296     node_.children = node_.children[:mid]
297     return node_.parent

```

```

291
292 def insert_node(_node):
293     log("inserting : {} in node: {} having children: {}".format(value,
294         ↳ _node.keys, getattr(_node, "children", "NULL")))
295     if _node.is_leaf: # logic for leaf node
296         log("node: {} is leaf".format(_node))
297         _index = bisect_right(_node.keys, value) # bisect and get index
298         ↳ value of where to insert value in node.keys
299         _node.keys.insert(_index, value)
300         if not _node.is_balanced:
301             split_leaf_node(_node)
302             log("----- Tree status after split-----")
303             log(self.__root)
304             log(self.__root.children)
305             log(_node.parent.children)
306             log(getattr(self.__root, "children", "NULL"))
307         else:
308             return
309     else: # logic for internal node
310         if not _node.is_balanced:
311             self.insert(split_internal_node(_node))
312         else:
313             _index = bisect_right(_node.keys, value)
314             log(_node.keys, _node.children, _index)
315             insert_node(_node.children[_index])
316
317 insert_node(self.__root)
318
319 @staticmethod
320 def traverse_left_to_right(node, index):
321     if node.children[index].is_leaf:
322         node.children[index + 1].keys.insert(0,
323             ↳ node.children[index].keys[-1])
324         node.children[index].keys.pop()
325         node.keys[index] = node.children[index + 1].keys[0]
326     else:
327         node.children[index + 1].children.insert(0,
328             ↳ node.children[index].children[-1])
329         node.children[index].children[-1].parent = node.children[index + 1]
330         node.children[index + 1].keys.insert(0, node.keys[index])
331         node.children[index].children.pop()
332         node.children[index].keys.pop()
333
334 @staticmethod
335 def traverse_right_to_left(node, index):
336     if node.children[index].is_leaf:
337         node.children[index].keys.append(node.children[index + 1].keys[0])

```



```
334         node.children[index + 1].keys.remove(node.children[index +
335         ↪ 1].keys[0])
336     node.keys[index] = node.children[index + 1].keys[0]
337 else:
338     node.children[index].children.append(node.children[index +
339     ↪ 1].children[0])
340     node.children[index + 1].children[0].parent = node.children[index]
341     node.children[index].keys.append(node.keys[index])
342     node.keys[index] = node.children[index + 1].children[0]
343     node.children[index + 1].children.remove(node.children[index +
344     ↪ 1].children[0])
345     node.children[index + 1].keys.remove(node.children[index +
346     ↪ 1].keys[0])
347
348 def delete(self, delete_value):
349     def merge(node, index):
350         if node.children[index].is_leaf:
351             node.children[index].keys = node.children[index].keys +
352             ↪ node.children[index + 1].keys
353             node.children[index].sibling = node.children[index + 1].sibling
354         else:
355             node.children[index].keys = node.children[index].keys +
356             ↪ [node.keys[index]] + node.children[
357             index + 1].keys
358             node.children[index].children = node.children[index].children +
359             ↪ node.children[index + 1].children
360             node.children.remove(node.children[index + 1])
361             node.keys.remove(node.children[index])
362             if node.keys:
363                 return node
364             else:
365                 node.children[0].parent = None
366                 self.__root = node.children[0]
367                 del node
368                 return self.__root
369
370 def delete_node(value, node):
371     log("deleting {} from node: {}".format(value, node))
372     if node.is_leaf:
373         log("node is leaf")
374         _index = bisect_left(node.keys, value)
375         try:
376             node_ = node.keys[_index]
377         except IndexError:
378             return False
379         else:
380             if node_ != value:
381                 return False
```

```

375         else:
376             node.keys.remove(value)
377             return True
378     else:
379         log("traversing internal node for deleting value")
380         _index = bisect_right(node.keys, value)
381         log("encountered index: {} having child values: {}".format(_index,
382             ↪ node.children[_index]))
382         if _index <= len(node.keys):
383             # print(node.children[_index].is_leaf)
384             # print(node.children[_index],
385             ↪ node.children[_index].total_keys,
386             ↪ node.children[_index].degree / 2,
387             ↪ node.children[_index].is_empty)
385         if not node.children[_index].is_empty:
386             return delete_node(value, node.children[_index])
387         elif not node.children[_index - 1].is_empty:
388             self.traverse_left_to_right(node, _index - 1)
389             return delete_node(value, node.children[_index])
390         else:
391             return delete_node(value, merge(node, _index))
392     delete_node(delete_value, self.__root)
393
394 def _test():
395     # test_lis = [0, 1, 11, 1, 2, 22, 13, 14, 4, 5, 23, 1, 51, 12, 31]
396     test_lis = [10, 1, 159, 200, 18, 90, 8, 17, 9]
397     # test_lis = range(10)
398     b = BPlusTree(degree=4)
399     for val in test_lis:
400         b.insert(val)
401         print("----- B+ TREE AFTER INSERT : {:3d} -----".format(val))
402         b.pretty_print()
403     # print("searching range.....")
404     # result = b.search_range(1, 12)
405     search_start, search_end = 1, 23
406     print("----- Searching in batch for {} to {} -----".format(search_start,
407         ↪ search_end))
407     print("Result*: {} \n (*distinct values)".format(b.search(search_start,
408         ↪ search_end)))
408     for delete_val in [200, 18, 9]:
409         print("----- DELETING {} -----".format(delete_val))
410         b.delete(delete_val)
411         print("----- B+ TREE AFTER DELETING : {:3d}
412             ↪ -----".format(delete_val))
412         b.pretty_print()
413
414
415 if __name__ == "__main__":

```

```

416 from collections import OrderedDict
417 choices = OrderedDict({
418     1: "Insert",
419     2: "Batch Insert",
420     3: "Delete",
421     4: "Search",
422     5: "Search Range",
423     6: "Terminate"
424 })
425 degree = input("Enter Degree of tree[4]: ")
426 b = BPlusTree(degree=int(degree) if degree else 4)
427 while True:
428     print("\n".join("{} {}".format(key, val) for key, val in
429         ↪ choices.items()))
429     choice = int(input("Enter Choice: "))
430     if choice == 1:
431         val = int(input("Enter number to insert: "))
432         b.insert(val)
433         print("----- B+ TREE AFTER INSERT : {:3d}
434         ↪ -----".format(val))
434         b.pretty_print()
435     elif choice == 2:
436         _values = map(int, input("Enter numbers (space separated):
437         ↪ ").split())
437         for val in _values:
438             b.insert(val)
439             print("----- B+ TREE AFTER INSERT : {:3d}
440             ↪ -----".format(val))
440             b.pretty_print()
441     elif choice == 3:
442         val = int(input("Enter number to delete: "))
443         b.delete(val)
444         print("----- B+ TREE AFTER DELETING : {:3d}
445         ↪ -----".format(val))
445         b.pretty_print()
446     elif choice == 4:
447         val = int(input("Enter number to search: "))
448         result = b.search(val, val)
449         print("Result*: {} \n (*distinct values)".format(b.search(result)))
450     elif choice == 5:
451         search_start = int(input("Enter start number of range: "))
452         search_end = int(input("Enter end number of range: "))
453         result = b.search(search_start, search_end)
454         print("----- Searching in batch for {} to {}
455         ↪ -----".format(search_start, search_end))
455         print("Result*: {} \n (*distinct
456         ↪ values)".format(b.search(search_start, search_end)))
456         # b.pretty_print()

```

```
457     else:
458         print("Thanks for using the service!!")
459         break
```

Output

```
Enter Degree of tree[4]: 4
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 1
Enter number to insert: 5
Leaf Node    : [5]    height >> 0
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice:
```

```

Leaf Node      : [5]      height >> 0
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 2
Enter numbers (space separated): 12 200 89 1
Leaf Node      : [5, 12]   height >> 0
Leaf Node      : [5, 12, 200] height >> 0
Internal Node  : [89]      height >> 0
Leaf Node      : [5, 12]   height >> 1
Leaf Node      : [89, 200] height >> 1
Internal Node  : [89]      height >> 0
Leaf Node      : [1, 5, 12] height >> 1
Leaf Node      : [89, 200] height >> 1
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice:

```

```

1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 1
Enter number to insert: 450
----- B+ TREE AFTER INSERT : 450 -----
Internal Node  : [50, 200, 320] height >> 0
Leaf Node      : [10, 25, 30]   height >> 1
Leaf Node      : [50, 110]      height >> 1
Leaf Node      : [200, 300]      height >> 1
Leaf Node      : [320, 400, 450] height >> 1
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate

```

```

1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 3
Enter number to delete: 320
----- B+ TREE AFTER DELETING : 320 -----
Internal Node : [50, 200, 320] height >> 0
Leaf Node      : [10, 25, 30] height >> 1
Leaf Node      : [50, 110] height >> 1
Leaf Node      : [200, 300] height >> 1
Leaf Node      : [400, 450] height >> 1
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 5
Enter start number of range: 10
Enter end number of range: 50
----- Searching in batch for 10 to 50 -----
Result*: [10, 25, 30, 50]

```

III. SUMMARY

- ✓ all leaves at the same lowest level
- ✓ all nodes at least half full (except root)
- ✓ Supports Range Query
- ✗ sequential search overhead may rise if large number of record in result of range query

Let f be the degree of tree and n be the total number of data then

Space Complexity				
	Max # pointers	Max # keys	Min # pointers	Min # keys
Non-leaf	f	$f - 1$	$\lceil f/2 \rceil$	$\lceil f/2 \rceil - 1$
Root	f	$f - 1$	2	1
Leaf	f	$f - 1$	$\lfloor f/2 \rfloor$	$\lfloor f/2 \rfloor$

- Number of disk accesses proportional to the height of the B+ tree. which is the *worst-case height* of a B+ tree is:

$$h \propto \log_f \frac{n+1}{2} \approx O(\log_f n) \quad (1)$$

Let f be the degree of tree and n be the total number of data then

Time Complexity		
	Time Complexity	Remarks
height	$O(\log_f n)$	
Root	$O(f \log_f n)$	linear search inside each nodes
search	$O(\log_2 f \log_f n)$	binary search inside each node
insert	$O(\log_f n)$	if splitting not require
insert	$O(f \log_f n)$	if splitting require
insert	$O(\log_f n)$	if merge not require
insert	$O(f \log_f n)$	if merge require