

Practical 3

- Implementation of B+ Tree Indexing for Database query processing
- Input-output for Select Query on exact Match, Range Query, Insert , delete Query.
- Analysis report

Implementation

```
#!/usr/bin/python3
# coding=utf-8
"""
some terminology of python:
_var => (convention only) underscore prefix is just a hint to programmer that a variable or method starting with a single underscore is intended for internal u
var_ => (convention only) to brake name conflict
__var => "dunders" (name mangling) rewrite the attribute name in order to avoid naming conflicts in subclasses.
        interpreter changes the name of the variable in a way that makes it harder to create collisions when the class is extended later.
"""

import math
import logging
import os
import random
from datetime import datetime
from bisect import bisect_right, bisect_left
from collections import deque

__author__ = "Gahan Saraiya"
DEBUG = False
LOG_DIR = "."
logger = logging.getLogger('bPlusTree')
logging.basicConfig(level=logging.DEBUG,
                    format='%(asctime)s [(name)-8s - %(levelname)s]: %(message)s',
                    datefmt='[%Y-%d-%m_%H.%M.%S]',
                    filename=os.path.join(LOG_DIR, 'b_plus_tree.log'),
                    filemode='w')

ch = logging.StreamHandler() # create console handler with a higher log level
ch.setLevel(logging.DEBUG)
# create formatter and add it to the handlers
formatter = logging.Formatter('%(asctime)s [(name)-8s - %(levelname)s]: %(message)s')
ch.setFormatter(formatter)
# add the handlers to the logger
logger.addHandler(ch)

def log(*msg):
    if DEBUG:
        logger.debug(msg)
    else:
        pass

class InternalNode(object):
    """
    Class : B+ Tree Internal Node
    represents internal (non-leaf) node in B+ tree
    """
    def __init__(self, degree=4):
        """
        initialize B tree node
        :param degree: specify degree of btree # default degree set to 4
        """
        self.degree = degree
        self.keys = [] # store keys/data values
        self.children = [] # store child nodes (list of instances of BTreeNode); empty list if node is leaf node
        self.parent = None

    def __repr__(self):
        return " | ".join(map(str, self.keys))

    @property
    def is_leaf(self):
        return False

    @property
    def total_keys(self):
        return len(self.keys)

    @property
    def is_balanced(self):
        # return False if total keys exceeds max accommodated keys (degree - 1)
        return self.total_keys <= self.degree - 1

    @property
    def is_full(self):
        return self.total_keys >= self.degree

    @property
```

```
def is_empty(self):
    return self.total_keys < (self.degree + 1) // 2

class LeafNode(object):
    """
    Class : B+ Tree Leaf Node
    represents leaf node in B+ tree
    """
    def __init__(self, degree=4):
        self.degree = degree
        self.keys = [] # data values
        self.sibling = None # sibling node to point
        self.parent = None # parent node - None for root node

    def __repr__(self):
        return " | ".join(map(str, self.keys))

    @property
    def is_leaf(self):
        return True

    @property
    def total_keys(self):
        return len(self.keys)

    @property
    def is_balanced(self):
        # return False if total keys exceeds max accommodated data (degree - 1)
        return self.total_keys < self.degree

    @property
    def is_full(self):
        return not self.is_balanced

    @property
    def is_empty(self):
        return self.total_keys < math.floor(self.degree / 2)

class BPlusTree(object):
    def __init__(self, degree=4):
        self.degree = degree
        self.__root = LeafNode(degree=degree)
        self.__leaf = self.__root

    def search_key(self, start_node, value):
        """
        :param start_node: get root or any non leaf node
        :param value: value to be search
        :return: most matching leaf node
        """
        if start_node.is_leaf:
            _index = bisect_left(start_node.keys, value)
            return _index, start_node
        else:
            _index = bisect_right(start_node.keys, value)
            return self.search_key(start_node.children[_index], value)

    def search(self, start=None, end=None):
        """
        :param start: specify start node to search range for
        :param end: specify end node for range search
        :return:
        """
        _result = []
        node = self.__root
        leaf = self.__leaf

        if start is None:
            while True:
                for value in leaf.keys:
                    if value <= end:
                        _result.append(value)
                    else:
                        return _result
                if leaf.sibling is None:
                    return _result
                else:
                    leaf = leaf.sibling
            elif end is None:
                _index, leaf = self.search_key(node, start)
                _result.extend(leaf.keys[_index:]) # equivalent to _result + leaf
                while True:
                    if leaf.sibling is None:
                        return _result
                    else:
                        leaf = leaf.sibling
                        _result.extend(leaf.keys)
            else:
                if start == end:
                    _index, _node = self.search_key(node, start)
                    try:
                        if _node.keys[_index] == start:
```

```

        _result.append(_node.keys[_index])
        return _result
    else:
        return _result
except IndexError:
    return _result
else:
    _index1, _node1 = self.search_key(node, start)
    _index2, _node2 = self.search_key(node, end)
    if _node1 is _node2:
        if _index1 == _index2:
            return _result
        else:
            _result.extend(_node1.keys[_index1:_index2])
            return _result
    else:
        _result.extend(_node1.keys[_index1:])
        node_ = _node1
        while True:
            if _node1.sibling == _node2:
                _result.extend(_node2.keys[:_index2 + 1])
                return _result
            else:
                try:
                    _result.extend(node_.sibling.keys)
                    node_ = node_.sibling
                except AttributeError:
                    return _result

def traverse(self, _node):
    _result = []
    _result.extend(_node.keys)
    if getattr(_node, "sibling", None) is None:
        return _result
    for i in range(0, len(_node.sibling))[:-1]:
        _result.extend(self.traverse(_node.sibling[i]))
    while True:
        pass

def pretty_print(self):
    # print("B+ Tree:")
    queue, height = deque(), 0
    queue.append([self.__root, height])
    while True:
        try:
            node, height_ = queue.popleft()
            # print("adding node: {}".format(node))
        except IndexError:
            return
        else:
            if not node.is_leaf:
                print("Internal Node : {:} \theight >> {}".format(node.keys, height_))
                if height_ == height:
                    height += 1
                queue.extend([[i, height] for i in node.children])
            else:
                print("Leaf Node      : {} \theight >> {}".format([i for i in node.keys], height_))

def insert(self, value):
    # log("parent:{} leaf:{} node:{}\tkeys:{}\t children:{}".format(node.parent, node.is_leaf, node, node.keys, getattr(node, 'children', '0')))

def split_leaf_node(node):
    log("splitting leaf node: {}".format(node.keys))
    mid = self.degree // 2 # integer division in python3
    new_leaf = LeafNode(self.degree)
    new_leaf.keys = node.keys[mid:]
    if node.parent is None: # None and 0 are to be treated as different value
        parent_node = InternalNode(self.degree) # create new parent for node
        parent_node.keys, parent_node.children = [node.keys[mid]], [node, new_leaf]
        node.parent = new_leaf.parent = parent_node
        self.__root = parent_node
    else:
        _index = node.parent.children.index(node)
        node.parent.keys.insert(_index, node.keys[mid])
        node.parent.children.insert(_index + 1, new_leaf)
        new_leaf.parent = node.parent
        if not node.parent.is_balanced:
            split_internal_node(node.parent)
    node.keys = node.keys[:mid]
    node.sibling = new_leaf
    log("{} --- {} --- {}".format(node, node.sibling, self.__root.children))

def split_internal_node(node_):
    mid = self.degree // 2 # integer division in python3
    new_node = InternalNode(self.degree)
    new_node.keys = node_.keys[mid:]
    new_node.children = node_.children[mid:]
    new_node.parent = node_.parent
    for child in new_node.children:
        child.parent = new_node # assign parent to every new child of current node
    if node_.parent is None: # again Note that None and 0 are not same but both treated as False in boolean
        # need to make new root if we are to split root node
        new_root = InternalNode(self.degree)
        new_root.keys = [node_.keys[mid - 1]]
        new_root.children = [node_, new_node]
        node_.parent = new_node.parent = new_root # set parent of newly created node

```

```
        self.__root = new_root # set new ROOT node
    else:
        # if node is not root internal node
        _index = node_.parent.children.index(node_)
        node_.parent.keys.insert(_index, node_.keys[mid - 1])
        node_.parent.children.insert(_index + 1, new_node)
        if not node_.parent.is_balanced:
            split_internal_node(node_.parent)
    node_.keys = node_.keys[:mid - 1]
    node_.children = node_.children[:mid]
    return node_.parent

def insert_node(_node):
    log("inserting : {} in node: {} having children: {}".format(value, _node.keys, getattr(_node, "children", "NULL")))
    if _node.is_leaf: # logic for leaf node
        log("node: {} is leaf".format(_node))
        _index = bisect_right(_node.keys, value) # bisect and get index value of where to insert value in node.keys
        _node.keys.insert(_index, value)
        if not _node.is_balanced:
            split_leaf_node(_node)
            log("----- Tree status after split-----")
            log(self.__root)
            log(self.__root.children)
            log(_node.parent.children)
            log(getattr(self.__root, "children", "NULL"))
        else:
            return
    else: # logic for internal node
        if not _node.is_balanced:
            self.insert(split_internal_node(_node))
        else:
            _index = bisect_right(_node.keys, value)
            log(_node.keys, _node.children, _index)
            insert_node(_node.children[_index])

insert_node(self.__root)

@staticmethod
def traverse_left_to_right(node, index):
    if node.children[index].is_leaf:
        node.children[index + 1].keys.insert(0, node.children[index].keys[-1])
        node.children[index].keys.pop()
        node.keys[index] = node.children[index + 1].keys[0]
    else:
        node.children[index + 1].children.insert(0, node.children[index].children[-1])
        node.children[index].children[-1].parent = node.children[index + 1]
        node.children[index + 1].keys.insert(0, node.keys[index])
        node.children[index].children.pop()
        node.children[index].keys.pop()

@staticmethod
def traverse_right_to_left(node, index):
    if node.children[index].is_leaf:
        node.children[index].keys.append(node.children[index + 1].keys[0])
        node.children[index + 1].keys.remove(node.children[index + 1].keys[0])
        node.keys[index] = node.children[index + 1].keys[0]
    else:
        node.children[index].children.append(node.children[index + 1].children[0])
        node.children[index + 1].children[0].parent = node.children[index]
        node.children[index].keys.append(node.keys[index])
        node.keys[index] = node.children[index + 1].children[0]
        node.children[index + 1].children.remove(node.children[index + 1].children[0])
        node.children[index + 1].keys.remove(node.children[index + 1].keys[0])

def delete(self, delete_value):
    def merge(node, index):
        if node.children[index].is_leaf:
            node.children[index].keys = node.children[index].keys + node.children[index + 1].keys
            node.children[index].sibling = node.children[index + 1].sibling
        else:
            node.children[index].keys = node.children[index].keys + [node.keys[index]] + node.children[
                index + 1].keys
            node.children[index].children = node.children[index].children + node.children[index + 1].children
    node.children.remove(node.children[index + 1])
    node.keys.remove(node.children[index])
    if node.keys:
        return node
    else:
        node.children[0].parent = None
        self.__root = node.children[0]
        del node
        return self.__root

def delete_node(value, node):
    log("deleting {} from node: {}".format(value, node))
    if node.is_leaf:
        log("node is leaf")
        _index = bisect_left(node.keys, value)
        try:
            node_ = node.keys[_index]
        except IndexError:
            return False
        else:
            if node_ != value:
                return False
            else:
```

```
        node.keys.remove(value)
        return True
    else:
        log("traversing internal node for deleting value")
        _index = bisect_right(node.keys, value)
        log("encountered index: {} having child values: {}".format(_index, node.children[_index]))
        if _index <= len(node.keys):
            # print(node.children[_index].is_leaf)
            # print(node.children[_index], node.children[_index].total_keys, node.children[_index].degree / 2, node.children[_index].is_empty)
            if not node.children[_index].is_empty:
                return delete_node(value, node.children[_index])
            elif not node.children[_index - 1].is_empty:
                self.traverse_left_to_right(node, _index - 1)
                return delete_node(value, node.children[_index])
            else:
                return delete_node(value, merge(node, _index))
        delete_node(delete_value, self.__root)

def _test():
    # test_lis = [0, 1, 11, 1, 2, 22, 13, 14, 4, 5, 23, 1, 51, 12, 31]
    test_lis = [10, 1, 159, 200, 18, 90, 8, 17, 9]
    # test_lis = range(10)
    b = BPlusTree(degree=4)
    for val in test_lis:
        b.insert(val)
        print("----- B+ TREE AFTER INSERT : {:3d} -----".format(val))
        b.pretty_print()
    # print("searching range.....")
    # result = b.search_range(1, 12)
    search_start, search_end = 1, 23
    print("----- Searching in batch for {} to {} -----".format(search_start, search_end))
    print("Result*: {} \n (*distinct values)".format(b.search(search_start, search_end)))
    for delete_val in [200, 18, 9]:
        print("----- DELETING {} -----".format(delete_val))
        b.delete(delete_val)
        print("----- B+ TREE AFTER DELETING : {:3d} -----".format(delete_val))
        b.pretty_print()

if __name__ == "__main__":
    from collections import OrderedDict
    choices = OrderedDict({
        1: "Insert",
        2: "Batch Insert",
        3: "Delete",
        4: "Search",
        5: "Search Range",
        6: "Terminate"
    })
    degree = input("Enter Degree of tree[4]: ")
    b = BPlusTree(degree=int(degree) if degree else 4)
    while True:
        print("\n".join("{} {}".format(key, val) for key, val in choices.items()))
        choice = int(input("Enter Choice: "))
        if choice == 1:
            val = int(input("Enter number to insert: "))
            b.insert(val)
            b.pretty_print()
        elif choice == 2:
            _values = map(int, input("Enter numbers (space separated): ").split())
            for val in _values:
                b.insert(val)
                b.pretty_print()
        elif choice == 3:
            val = int(input("Enter number to delete: "))
            b.delete(val)
            b.pretty_print()
        elif choice == 4:
            val = int(input("Enter number to search: "))
            b.search(val, val)
            b.pretty_print()
        elif choice == 5:
            start = int(input("Enter start number of range: "))
            end = int(input("Enter end number of range: "))
            b.search(start, end)
            b.pretty_print()
        else:
            print("Thanks for using the service!!")
            break
```

Results

```
Enter Degree of tree[4]: 4
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 1
Enter number to insert: 5
Leaf Node      : [5]      height >> 0
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice:
```

```
Leaf Node      : [5]      height >> 0
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 2
Enter numbers (space separated): 12 200 89 1
Leaf Node      : [5, 12]   height >> 0
Leaf Node      : [5, 12, 200] height >> 0
Internal Node : [89]      height >> 0
Leaf Node      : [5, 12]   height >> 1
Leaf Node      : [89, 200] height >> 1
Internal Node : [89]      height >> 0
Leaf Node      : [1, 5, 12] height >> 1
Leaf Node      : [89, 200] height >> 1
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice:
```

```
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 1
Enter number to insert: 450
----- B+ TREE AFTER INSERT : 450 -----
Internal Node : [50, 200, 320] height >> 0
Leaf Node      : [10, 25, 30] height >> 1
Leaf Node      : [50, 110] height >> 1
Leaf Node      : [200, 300] height >> 1
Leaf Node      : [320, 400, 450] height >> 1
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
```

```
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 3
Enter number to delete: 320
----- B+ TREE AFTER DELETING : 320 -----
Internal Node : [50, 200, 320] height >> 0
Leaf Node      : [10, 25, 30] height >> 1
Leaf Node      : [50, 110] height >> 1
Leaf Node      : [200, 300] height >> 1
Leaf Node      : [400, 450] height >> 1
1 Insert
2 Batch Insert
3 Delete
4 Search
5 Search Range
6 Terminate
Enter Choice: 5
Enter start number of range: 10
Enter end number of range: 50
----- Searching in batch for 10 to 50 -----
Result*: [10, 25, 30, 50]
```

Analysis

- all leaves at the same lowest level
 - all nodes at least half full (except root)
- Let f be the degree of tree and n be the total number of data then

	Max # pointers	Max # keys	Min # pointers	Min # keys
Non-leaf	f	$f - 1$	$\lceil f/2 \rceil$	$\lceil f/2 \rceil - 1$
Root	f	$f - 1$	2	1
Leaf	f	$f - 1$	$\lfloor f/2 \rfloor$	$\lfloor f/2 \rfloor$

- Number of disk accesses proportional to the height of the B-tree.
- The **worst-case height** of a B+ tree is:

$$h \propto \log_f \frac{n+1}{2} \sim O(\log_f n)$$

	Time Complexity	Remarks
height	$O(\log_f n)$	
search	$O(f \log_f n)$	linear search inside each nodes
search	$O(\log_2 f \log_f n)$	binary search inside each node
insert	$O(\log_f n)$	if splitting not require
insert	$O(f \log_f n)$	if splitting require
insert	$O(\log_f n)$	if merge not require
insert	$O(f \log_f n)$	if merge require